


Random Forests:

We used random samples (rows) and random features as well (columns). We first split on 128 which was giving a good enough accuracy of 67, but we wanted to improve the accuracy, so we checked each value on each column and associated gini index and the one with the lowest gini index is selected. So it assures that we are taking the best value, instead of taking 128 as the split point (fixed threshold). Also the max depth of the tree we went till 40. When we took it to 80, there was no significant improvement in accuracy, and hence we set it to 40 and the min leaf we set to 1 even though it has a chance of overfitting because it gives a good accuracy. We created 20 decision trees randomly and we took a vote among all of them to classify each image.

##	Number.of.trees	Training.time	Testing.time	Accuracy
## 1	5	1m 1s	0.318s	71.26
## 2	10	2m 8s	0.383s	72.11
## 3	15	3m 10s	0.453s	63.73
## 4	20	4m 8s	0.514s	72.11
## 5	25	5m 11s	0.577s	79.11
## 6	30	5m 59s	0.650s	87.91
## 7	35	6m 58s	0.712s	79.64

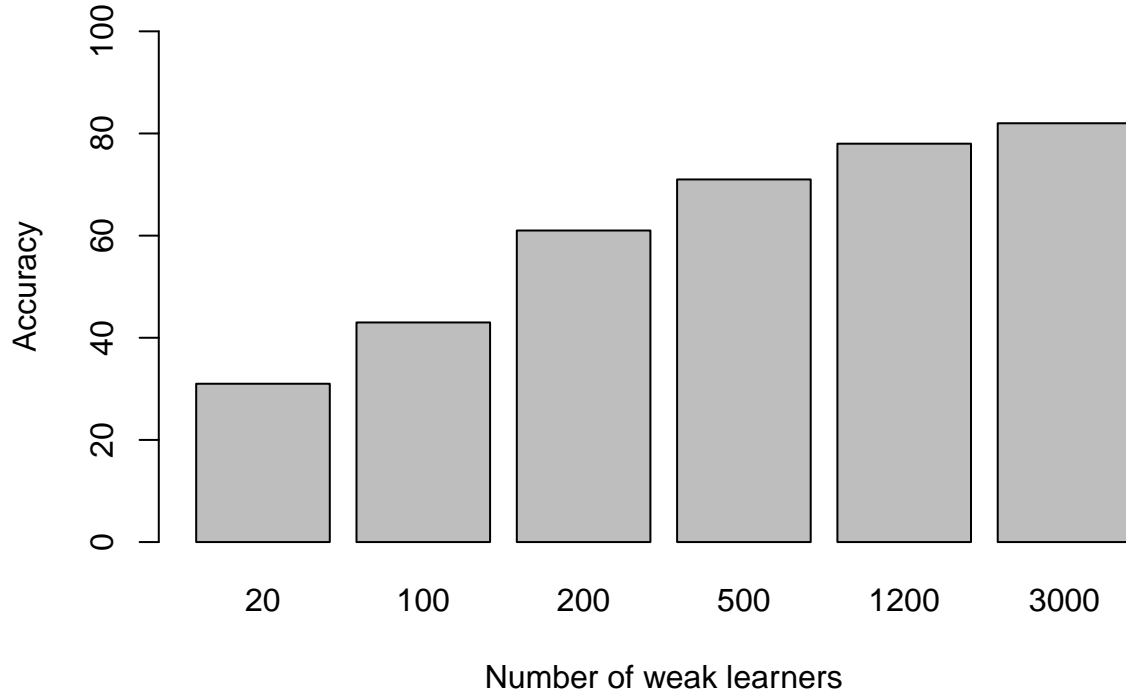


AdaBoost:

In Adaboost, we used pair of columns as a stub where our stub (weak learner) hypothesizes as subtraction of values at the given column pairs, and if it is less than 0, we label it as positive, otherwise negative. Since Adaboost works only for binary classification, we had to improvise to make 4 one-versus-all classifiers, and using these classifiers and the learned parameters a_i , we calculated the H values for a single image in 4 classifiers and the one with highest positive value was given its respective corresponding class label. If none of the classifiers were able to classify, then we give it a random class. We started with 20 learners giving an accuracy of 31%, going up to 1,200 learners giving an accuracy of 78% but incorporating more learners ~3000 didn't provide significant increase in accuracy. And also considering the time trade-off, we stopped at 3,000 learners.

##	Number.of.weak.learners	Training.time	Testing.time	Accuracy
----	-------------------------	---------------	--------------	----------


## 1	20	10s	57s	31
## 2	100	1m	1m 3s	43
## 3	200	2m 27s	1m 11s	61
## 4	500	3m 42s	1m 19s	71
## 5	1200	5m 3s	1m 24s	78
## 6	3000	16m 11s	1m 30s	82



K-Nearest Neighbours:

To find the nearest neighbours, we used the Euclidean distance. We analyzed the accuracy for different values of K and found out the best performance at K=11. Though K=43 gives a slightly better accuracy, the extra time is not worth the increase in accuracy.

##	K	Testing.time	Accuracy
## 1	1	5m 53s	67.23
## 2	2	6m	67.23
## 3	3	6m 18s	68.72
## 4	4	6m 14s	69.67
## 5	5	6m 22s	69.88
## 6	6	6m 41s	69.77
## 7	8	6m 34s	69.35
## 8	10	6m 34s	70.41
## 9	11	6m 52s	71.26
## 10	12	6m 5s	70.51
## 11	14	6m 21s	69.88
## 12	39	7m 32s	71.36
## 13	43	7m 47s	71.68



Among all classifiers, taking time and computation into consideration, we would suggest AdaBoost to a client even though random forest gives high accuracy in some cases but the accuracies are erratic and are subject to a lot of parameter tuning.